

Not Too Early, Not All at Once: Design Tensions in AI-Mediated Self-Disclosure in Online Dating

Pei-Hua Tsai
National Yang Ming Chiao Tung
University
Department of Computer Science
Hsinchu, Taiwan
phtsai.cs11@nycu.edu.tw

Tianyi Zhang
Singapore Management University
School of Computing and Information
Systems
Singapore, Singapore
tianyizhang.2023@smu.edu.sg

Emran Poh
Singapore Management University
School of Computing and Information
Systems
Singapore, Singapore
emran.poh.2025@smu.edu.sg

Anthony Tang
Singapore Management University
School of Computing and Information
Systems
Singapore, Singapore
tonyt@smu.edu.sg

Yung-Ju Chang
National Yang Ming Chiao Tung
University
Department of Computer Science
Hsinchu, Taiwan
armuro@nycu.edu.tw

Abstract

Online dating relies on self-disclosure, yet initial conversations are fragile: users must navigate uncertainty around timing, boundaries, and reciprocity with little shared context. While advances in AI raise the possibility of mediating disclosure, how such support might reshape the experience of early-stage relational disclosure remains underexplored. We conducted 29 semi-structured interviews to examine how daters envision AI-mediated self-disclosure in online dating. Our findings surface recurring design tensions rather than simple opportunities or risks. Participants welcomed guidance that could pace disclosure, support reflection, and reduce social awkwardness, but stressed preserving agency and authorship. They valued interpretive assistance for sense-making of ambiguous partner cues, yet worried that algorithmic interpretation might foreclose gradual discovery. Participants also described relational buffering as face-saving, while cautioning that increased efficiency risks undermining reciprocity, surprise, and co-constructed intimacy. Together, these findings suggest that designing AI for intimate contexts requires attending to how support redistributes agency, interpretation, and participation over time, rather than treating disclosure as an optimization problem.

CCS Concepts

• **Human-centered computing** → **Empirical studies in HCI**; **Collaborative and social computing**; *Interaction design theory, concepts and paradigms.*

Keywords

AI-mediated communication; Self-disclosure; Online dating; Human-AI interaction; Relationship development

ACM Reference Format:

Pei-Hua Tsai, Tianyi Zhang, Emran Poh, Anthony Tang, and Yung-Ju Chang. 2018. Not Too Early, Not All at Once: Design Tensions in AI-Mediated Self-Disclosure in Online Dating. In *Proceedings of Make sure to enter the correct conference title from your rights confirmation email (Conference acronym 'XX)*. ACM, New York, NY, USA, 12 pages. <https://doi.org/XXXXXXX.XXXXXXX>

1 Introduction

Online dating has become one of the most common ways relationships begin [40]. In the U.S., around 30% of adults report having used a dating app or site, and among 18–29 year-olds that number rises to roughly half [35]. Yet initiating connections in this context presents distinct challenges. Conversations often stretch across days or weeks through asynchronous, text-based exchanges, where activity that spikes and stalls unpredictably, making momentum hard to sustain [32, 55]. The absence of immediate feedback limits users' ability to calibrate their approach or repair missteps in real time [18, 32, 55]. At the same time, users know little about a partner's preferences, boundaries, or communication style, yet feel pressure to present the best version of themselves and establish rapport within a short window after matching [17, 32, 39, 55]. These challenges are further amplified in cross-cultural encounters, where differing norms and expectations complicate judgments about what to say and how to respond [5, 42, 48]. Together, the asynchronous rhythm of exchanges, the lack of shared physical context, and the compressed relational timelines of dating platforms leave users especially uncertain about the timing, amount, and style of disclosure [14, 32, 55].

We use *self-disclosure*—the intentional sharing of personal information—to build intimacy [13]. Disclosure is fundamental to relationship development, yet in online dating it must be continually modulated under conditions of uncertainty [18, 32, 55], time pressure [32, 39], and impression management [14, 17, 51, 55]. Missteps—such as

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, to republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.
Conference acronym 'XX, Woodstock, NY

© 2018 Copyright held by the owner/author(s). Publication rights licensed to ACM.
ACM ISBN 978-1-4503-XXXX-X/2018/06
<https://doi.org/XXXXXXX.XXXXXXX>

disclosing too much too soon, or letting participation become one-sided, or failing to align with a partner's cues—can quickly disrupt connection [12, 47].

Recent advances in AI create new opportunities for addressing these disclosure challenges. Even with limited shared context, AI can analyze conversational context and generate tailored suggestions [33]. For example, by interpreting conversational flow [43], AI systems could help users pace disclosures appropriately; by monitoring participation patterns [45], they could surface risks to reciprocity when conversations become one-sided; and by drawing on contextual cues [52], they could help users attune messages to a partner's emerging communication style. Similarly, we have seen that AI can scaffold communication in dyadic interactions. AI can suggest topics of shared interest to initiate conversation [29, 45], support message writing during relationship dissolution [15], or assist with tone adjustment [43].

However, how such support reshapes the experience of early-stage relational disclosure, where ambiguity and gradual mutual discovery are central, remains underexplored. In particular, little is known about how AI-mediated support redistributes who decides what is said, how partners are understood, and how users remain involved in the unfolding interaction. This research gap is especially important for understanding how AI can support self-disclosure in online dating, where fragile relational dynamics shape not only what is shared, but how meaning, trust, and connection are gradually formed.

To address this gap, we examine how users reason about both the potential benefits and the risks of AI agents mediating self-disclosure in online dating. Specifically, we ask:

- (1) How do users envision AI agents supporting self-disclosure in early-stage online dating interactions?
- (2) What concerns do users raise about how such AI-mediated support might reshape the experience of self-disclosure?

To address these questions, we conducted 29 semi-structured interviews with individuals who had prior experience using online dating apps. Rather than evaluating a specific prototype, we adopted an exploratory design approach. At the outset, participants were briefly shown illustrative provocations—speculative materials inspired by critical design probes [6, 16]—that depicted possible forms of AI mediation of self-disclosure. These provocations were used to spark imagination about how such mediation might unfold in dating contexts. Participants were then invited to recall or imagine situations where self-disclosure felt mistimed, uneven, or socially risky, and to consider how AI agents might have intervened to help or complicate matters.

Our analysis shows that participants saw AI agents as potentially helpful in pacing conversations, easing delicate disclosures, and attuning sharing to a partner's cues. At the same time, they raised concerns about how such mediation might change the experience of disclosure itself, particularly around authorship, interpretation, and involvement in the interaction. Together, these reflections point to tradeoffs that shape when AI support feels helpful and when it feels intrusive in early-stage relational disclosure.

This paper contributes to HCI research on AI-mediated self-disclosure in online dating in three ways:

- First, we surface three recurring design tensions—*Disclosure Agency*, *Interpretation*, and *Relational Buffering*—that characterize how AI mediation in early dating redistributes authorship, sensemaking, and participation in self-disclosure.
- Second, we build on prior work on self-disclosure in relationship development by contributing an empirical account of how users envision AI agents scaffolding, delegating, and interpreting disclosure in the fragile early stages of online dating—expectations distinct from prior work on topic suggestion or message assistance.
- Third, we introduce the idea of multi-agent mediation—where each partner is supported by an agent that coordinates across perspectives—as a conceptual direction for future AI design in relational settings.

2 Related Work

2.1 Self-Disclosure in Romantic Relationships

Self-disclosure is a core practice in human relationships [13], and involves the intentional sharing of personal thoughts, feelings, or experiences. It serves different purposes across contexts: cultivating closeness in families and friendships [10, 28], establishing credibility in workplaces [36, 46], reducing social distance in classrooms [9, 19], or fostering trust in therapy [22, 24]. In each setting, norms and expectations shape disclosure, requiring continual choices about what, when, and to whom to reveal.

Romantic relationships highlight these dynamics in especially consequential ways. Social Penetration Theory depicts intimacy as a gradual process of peeling back layers of the self, with closeness developing through reciprocal disclosures over time [2]. Yet disclosure is never risk-free. Each revelation can threaten face and unsettle impressions, as Politeness Theory emphasizes [8]. Reciprocity is a fragile norm: uneven participation, where one partner shares more than the other, can hinder mutual discovery and weaken intimacy [10]. Appropriateness and pacing also matter, since disclosures that come too early, too late, or at the wrong depth can disrupt rapport and stall relational progress [12, 20]. Attunement to a partner's engagement is equally significant. The Interpersonal Process Model of Intimacy shows that intimacy develops when disclosures are met with validating and responsive reactions [30, 44]; when disclosures misalign with a partner's interest or responsiveness, they may generate discomfort, reduce liking, or prematurely end the interaction [47]. Taken together, these perspectives illustrate disclosure as a delicate balancing act: intimacy develops through reciprocity, pacing, engagement, and attunement, but every act of sharing carries risks of misalignment, face loss, or relational withdrawal.

Online dating environments amplify these difficulties. Prior work highlights how reduced cues make it hard to assess whether disclosures are welcomed, overlooked, or misunderstood [14, 18]. Walther's Hyperpersonal Model adds that such environments can paradoxically accelerate intimacy through selective self-presentation and idealized partner perceptions, while simultaneously heightening risks of misinterpretation [49]. Cultural variation further complicates disclosure decisions: studies of relational mobility and sexual disclosure norms show that what counts as appropriate

disclosure differs significantly across contexts [42, 48]. These uncertainties make impression management especially salient, as daters must carefully balance authenticity with strategic self-presentation to appear appealing without oversharing or seeming disingenuous [14, 17, 18, 55]. Beyond ambiguity, platform design also compresses the timeline of relationship initiation. Rosen et al. [39] and Rosenfeld et al. [40] emphasize how online dating accelerates relational progress, creating pressure to reveal more in less time, while LeFebvre [32] and Satyaningrum et al. [41] highlight how this speed often produces premature intimacy, uneven pacing, or abrupt disconnection.

Taken together, prior work characterizes self-disclosure in online dating as high-stakes, shaped by uncertainty, fragile reciprocity, and relational risk. Users must judge how much to share and when, interpret sparse or delayed feedback, and manage impressions under compressed timelines, all while engaging a partner whose preferences and intentions remain largely unknown. These dynamics provide the context for examining how AI might shape the way people navigate self-disclosure in online dating.

2.2 AI Mediation for Dyadic Interaction

Research on AI-mediated communication (AI-MC) has increasingly examined how computational systems shape disclosure in mediated interactions. Hancock et al. define AI-MC as communication in which a computational agent modifies, augments, or generates messages on behalf of a communicator to accomplish interpersonal goals [21]. Building on this definition, subsequent studies have investigated how AI can scaffold dyadic interaction.

One body of research demonstrates how AI can support conversational momentum and participation. Shin et al.'s IntroBot detects when conversations stall and suggests shared-interest topics, encouraging quieter participants to contribute [45]. Kim et al. show that keyword-based message expansion can facilitate continuous texting in romantic encounters, reducing awkward silences and stalled exchanges [29]. Sharma et al. demonstrate that real-time feedback from an AI collaborator enables peers to express themselves with greater empathy and confidence in text-based mental health support [43].

Other work emphasizes personalization and adaptive support in AI-mediated interaction. Zhou et al.'s account of XiaoIce demonstrates how empathetic chatbots can track users' moods and histories to calibrate responses dynamically, creating the sense of a more reciprocal relationship [52]. Li et al. propose benchmarks for personalized conversational modeling that point toward systems capable of tailoring interactional support to individuals' preferences and styles [33]. These efforts frame AI mediation as increasingly attuned to relational nuance rather than generic assistance.

AI has also been explored as a means of reducing interpersonal risk in sensitive exchanges. Fu et al. document how people imagine relying on AI to assist with breakup communication [15]. Hohenstein and Jung describe how AI systems act as a "moral crumple zone," absorbing blame when communication misfires and thereby reducing the interpersonal costs of poorly received messages [25]. In dating contexts, Zytke and Furlo's participatory design work envisions consent technologies that mediate sexual communication across online and offline interactions, normalizing ongoing

dialogue about boundaries and expectations rather than one-time consent events [53]. Relatedly, Zytke et al. design AI agents for social VR dating that negotiate or enforce consent on users' behalf, actively mediating boundary-setting in real time [56].

At the same time, a growing body of work cautions that AI mediation can introduce new relational costs. In dating contexts, Wu and Kelly show that dating profiles perceived as AI-generated are judged less trustworthy, even when attractiveness ratings remain unaffected [50]. Jakesch et al. further demonstrate that even the suspicion of AI involvement can undermine perceived credibility in online interactions [26]. Barkallah et al. extend these concerns by showing how users weigh the convenience of automation against risks of misrepresentation and diminished authenticity in AI-generated dating profiles [3]. These findings suggest that AI assistance, while helpful, can disrupt how sincerity and responsibility are attributed in relational settings.

Concerns also extend to how algorithmic mediation shapes interpretation and expression. Scholars warn that AI-generated suggestions may homogenize communication styles, nudging users toward dominant cultural norms and narrowing the diversity of self-expression [1]. Datey and Zytke further show that while users prefer personalized models for detecting online harm, they remain concerned about amplifying subjective judgments or reinforcing problematic stereotypes [11].

Taken together, prior work presents AI mediation as both enabling and consequential for dyadic interaction. While AI systems can scaffold participation, personalize support, and absorb social risk, they may also reshape authorship, interpretation, and trust. Importantly, little research has examined how users themselves make sense of and navigate these tradeoffs in the context of early romantic self-disclosure—where ambiguity, gradual discovery, and relational vulnerability are central. This gap motivates our investigation into how people envision AI agents mediating self-disclosure in online dating, and how they negotiate when such support feels helpful versus intrusive.

3 Methods

We conducted an interview-based study to examine how users reason about both the potential benefits and the risks of AI agents mediating self-disclosure in online dating. Specifically, we investigated (1) how users envision AI agents supporting self-disclosure in early-stage online dating interactions, and (2) what concerns they raise about how such AI-mediated support might reshape the experience of self-disclosure. To explore these questions, we designed an interview-based study where participants were first exposed to the idea of AI mediation with a set of illustrative provocations. Participants were then asked to recall past awkward or past exchanges in online dating apps.¹ They then used these grounded recollections to speculate on the ways in which AI mediation might have been able to help (or hinder) those past experiences. By anchoring speculative reasoning in recalled interactions, participants could reason about AI mediation not as abstract futures, but as interventions into concrete relational breakdowns they had experienced.

¹Participants were asked only to consider awkward conversations, and not to disclose intimate or traumatic experiences, so as to minimize potential discomfort during the study. Participants were asked only to recount it from their perspective, and not to show us the conversation.




Figure 1: AI mediation provocations used as design probes, illustrated through the pacing scenario. Here, the participant is sharing personal information too quickly early in the interaction. Provocations were organized along two dimensions: *Perspective* and *Delivery Mode*. Under *Perspective*, the agent intervention either focused on *User Behavior*, highlighting what the participant had done in the interaction (e.g., “You’ve shared a lot of personal detail very quickly”), or on *Partner Experience*, highlighting how the partner might experience the interaction (e.g., “Your partner may feel overwhelmed by the pace of disclosure”). Under *Delivery Mode*, the agent interventions were either *Public*, visible to both the participant and partner, or *Private*, visible only to the participant.

3.1 AI Mediation Provocations

We prepared a set of illustrative provocations, inspired by critical design probes and design fiction [6, 16], using open-ended cues to stimulate speculation about possible forms of AI agent involvement. These provocations were not treated as features to be evaluated or as an analytical lens, but as open-ended probes to elicit reasoning about disclosure timing, responsibility, and interpretation.

Provocations were structured around two scenarios that are relevant in online dating and have been shown to shape relational development [12, 20, 32, 41, 47]: (1) *pacing*, or the timing and appropriateness of disclosures relative to the stage of the relationship (e.g., sharing too much too early or waiting too long to disclose), and (2) *reciprocity*, or the balance of sharing between participants (e.g., dominating the conversation or contributing very little). These scenarios guided participants in recalling past conversations and imagining how AI could support or influence those interactions.

Within each scenario, illustrative agent interventions (short, hypothetical system messages shown to participants) varied along

two dimensions, forming a 2×2 structure. The first dimension, *delivery mode*, distinguished between *private* agent prompts shown only to the participant and *public* agent prompts shown to both partners. The second dimension, *perspective*, differentiated whether the agent prompt framed the interaction in terms of the participant’s own *behavior* (i.e., what the participant had done in the conversation) or the partner’s *experience* (i.e., how the interaction might be felt or interpreted by the partner). These dimensions were refined through multiple iterative pilot studies, which focused on ensuring the provocations were clear, non-directive, and capable of eliciting divergent reactions rather than convergent judgments.

Figure 1 illustrates this 2×2 structure using the pacing scenario, with the reciprocity scenario provided in the Appendix (Figure 3).

3.2 Interview Protocol

At the start of each session, participants were informed of the study’s purpose and provided written consent.

Phase 1: AI Mediation Provocations. Participants were first introduced to the AI mediation provocations, which illustrated how AI might operate in dating apps. These served only as starting points; participants were reminded that AI mediation could extend beyond the examples shown. This phase lasted about 10 minutes.

Phase 2: Recalling Awkward Past Exchanges & Speculating on Design. Next, participants recalled dating-app conversations that felt awkward, unbalanced, or emotionally difficult. They described the situation, why it was challenging, and how they responded. While most examples came from participants’ own experiences, the interviewer occasionally prompted speculation about additional scenarios if relevant aspects had not been addressed. The interviewer used the two interaction scenarios—*pacing* and *reciprocity*—to guide reflections on the timing and balance of disclosures and alignment with partner cues.

For each recalled interaction, participants were first asked to envision how AI agents might have supported, mediated, or complicated the exchange based on their own ideas. After this initial reflection, the interviewer used the previously introduced AI mediation provocations to guide discussion, organized along a 2×2 structure defined by *delivery mode* (private messages visible only to the participant vs. public messages visible to both partners) and *perspective* (framing the situation in terms of the participant’s own *behavior* vs. the partner’s *experience*). We encouraged participants to consider additional possibilities beyond those mentioned, and to compare conditions, discussing differences in how they imagined the AI would act, how they might respond, and how the interaction could be affected under each combination.

All interviews were video-recorded, transcribed verbatim, and anonymized. Sessions lasted 90–120 minutes, and participants were compensated for their time. The protocol was approved by our institutional IRB.

3.3 Participants

We recruited 29 individuals who had used dating apps within the past year. Participants were recruited via online sign-up forms circulated through social media and community channels. As part of the form, we collected demographic information (age, gender,

ID	Age	Gen.	Education	Serious Rs.	Dating Apps Used	Usage Frequency
P01	24	F	Master	5	Pikabu	Several times/week
P02	24	M	Master	5	Jimu, Soul, Tantan, Tinder	Several times/day
P03	27	M	Master	5	Soul, TanTan, Tinder	<1/month
P04	27	NB	Master	2	Bumble, Feeld, Hinge, Tinder	Several times/day
P05	28	M	Master	4	Bumble, Feeld, Hinge, Tinder	Several times/week
P06	24	M	Bachelor	5	Bumble, Coffee Meets Bagel, Tinder	Several times/day
P07	23	F	Bachelor	0	Bumble, Hinge	Several times/day
P08	28	M	Master	4	Bumble, Coffee Meets Bagel, Hinge, Tinder	Several times/week
P09	21	F	High school	2	Bumble, Hinge	Several times/week
P10	29	F	Master	1	Bumble, Coffee Meets Bagel, Hinge	Several times/day
P11	29	M	Bachelor	2	Bumble, Tinder	Several times/month
P12	25	M	Bachelor	2	Coffee Meets Bagel	Several times/week
P13	27	M	Bachelor	2	Omi, Pikabu, Tinder	Several times/week
P14	25	M	Master	0	Coffee Meets Bagel, Tinder	Several times/week
P15	25	F	Bachelor	2	Bumble, Coffee Meets Bagel, Tinder	Several times/month
P16	25	F	Master	–	Bumble, Coffee Meets Bagel, Goodnight, Omi, Pikabu, Tinder	Several times/day
P17	25	F	Master	4	Coffee Meets Bagel, Pikabu	<1/month
P18	28	M	Master	1	Coffee Meets Bagel, Pikabu	Several times/day
P19	30	F	Master	2	Bumble, Pikabu, Soga	Several times/week
P20	23	F	Bachelor	2	Coffee Meets Bagel, Omi, Pikabu	Several times/week
P21	26	M	Master	2	Bumble, Coffee Meets Bagel, Pikabu	Several times/week
P22	23	F	Bachelor	2	Bumble, Coffee Meets Bagel, Tinder	Several times/day
P23	22	F	Master	5	Bumble, Coffee Meets Bagel, SOGA, Tinder	Several times/week
P24	40	F	Bachelor	2	Omi, Tinder	Several times/week
P25	31	M	Master	2	Bumble, Coffee Meets Bagel, Omi, Tinder	Several times/week
P26	46	M	Master	2	Omi, PGTalk, Tinder	Several times/month
P27	43	F	Bachelor	6	Coffee Meets Bagel	Several times/week
P28	35	M	Master	10	Coffee Meets Bagel, Gsland, Hornet, Jack'd, Omi, SCRUFF, Tinder	Several times/day
P29	33	F	Master	5	Bumble, Coffee Meets Bagel, Pikabu, Tinder	Several times/day

Table 1: Participant demographics and dating app usage patterns (N=29). Gen.: M=Male, F=Female, NB=Non-binary; Serious Rs.=Serious Relationships.

highest completed education, number of serious romantic relationships—defined as committed partnerships rather than casual or purely sexual encounters) and dating app usage (apps used and frequency), which we used to ensure diversity in habits and relationship histories. Table 1 summarizes the sample: 14 men, 14 women, and one non-binary participant, ranging in age from 21 to 46. Reported serious relationships ranged from 0 to 10, and participants described varied app practices, including use of dating platforms such as Tinder, Bumble, Coffee Meets Bagel, Pikabu, and Omi, with frequencies ranging from several times per day to less than once per month.

3.4 Data Analysis

Our data analysis and data collection proceeded in parallel. The analysis was informed by an inductive thematic analysis approach [7], in which the codebook was developed iteratively. We used MAXQDA² to manage, organize, and trace coding decisions throughout the process.

In our analysis, we examined how participants described challenges in self-disclosure during online dating interactions and how

they reasoned about AI-mediated support in addressing those challenges. We focused on participants' reflections on which aspects of the interaction required understanding or reflection, and how AI support might intervene in those moments, including how the visibility and framing of AI involvement could shape interactional dynamics. We also attended to participants' concerns about potential breakdowns introduced by AI mediation, such as risks related to agency, misinterpretation, and disruptions to relational processes, as well as the perceived consequences for participants themselves, their perceptions of their partners, and the developing relationship. Design tensions were identified by examining moments where participants articulated both the appeal and the risk of the same form of AI support, either within a single account or across contrasting accounts.

An initial set of codes was generated from one researcher's notes on the first five transcripts, accompanied by analytic memos that captured the meanings of emerging codes and categories. After this preliminary coding, three members of the research team jointly examined the codebook to compare codes and memos and refine the developing categories. As additional data were collected, one researcher coded every five transcripts and then met with the other two team members to review and refine the coding. New codes

²<https://atlasti.com>

were added, revised, or collapsed as needed, and their applicability was checked across earlier transcripts to ensure coherence. These iterative discussions not only supported consensus on the thematic structure but also fostered reflexivity in interpreting the data. This process ultimately yielded a set of themes capturing (1) aspects of interactions requiring understanding or reflection, (2) how AI support might intervene in those moments, including its visibility and framing, and (3) the perceived consequences of such interventions, including their underlying reasons and their impact on the self, the partner, and the developing relationship.

4 Findings: Design Tensions in AI-Mediated Self-Disclosure

Participants described a wide range of ways AI might support early dating conversations, particularly in moments of uncertainty, vulnerability, or social risk. At the same time, their accounts revealed recurring concerns about how such support might reshape the process of self-disclosure itself. Rather than centering on specific features, participants often focused on questions about control, interpretation, and participation in the interaction. We organize these concerns into three design tensions that highlight tradeoffs in how AI-mediated support redistributes decision-making, sensemaking, and presence in early relational encounters.

4.1 Disclosure Agency: Scaffolding Disclosure vs Preserving Authorship and Responsibility

Early dating conversations often unfold in bursts of instinctive messaging, where users type quickly, react emotionally, and only later realize they may have overshared, dominated the exchange, or mistimed their disclosures. Participants described moments where they were “typing too fast and not thinking about the situation” (P11), sending long messages without realizing how they might land, or continuing to talk past the moment when interest or momentum had peaked (P12). Because feedback is delayed or ambiguous, these misalignments often become apparent only after the fact, leaving users uncertain whether they strengthened connection or undermined it.

Participants saw AI agents as potentially helpful in interrupting these instinct-driven moments by heightening awareness without taking control. P11 emphasized how real-time support could slow impulsive disclosure:

“A lot of times [in dating] you just go by instincts and just type it out. That’s where the agent can help. [...] It could suggest something while I’m typing so fast and not thinking about the situation. So I think it’s about instantaneous suggestions—based on how I’m typing, what I’m typing.”

Others described pacing not only within messages but across turns. As P12 explained, knowing when to stop could preserve anticipation rather than exhaust it:

“I think in chatting, you shouldn’t just keep going endlessly. It’s best to stop when the mood is at its peak, because then there will be excitement for the next time. [...] So I think the agent could remind me when to wrap up, helping me stop at the best moment.”

More broadly, participants wanted AI support that heightened awareness without taking control. They envisioned agents nudging them in real time: slowing them down when they typed impulsively, encouraging participation when they stayed quiet, or flagging when message length or timing might overwhelm a partner. Crucially, guidance needed to be concrete and actionable rather than evaluative or vague. As P01 reflected, *“I might not realize that this kind of message length makes the other person uncomfortable, since it’s just the way I’m used to talking.”* When feedback lacked specificity, it failed to help: *“If I only got a reminder like ‘you talk too much,’ I wouldn’t know what exactly that means or how to fix it.”*

Beyond moment-to-moment nudges, participants also welcomed reflective coaching across conversations. They imagined reviewing past interactions to recognize recurring patterns and learn from mistakes. As P18 put it, such feedback could help them *“use that feedback to learn—so next time I’d know how to respond differently or share in a better way.”*

Yet several participants cautioned that the same forms of guidance they found helpful could become problematic if sustained or over-applied. Some worried that reliance on AI might “standardize” interaction styles, producing formulaic conversations that eroded individuality. As P20 reflected:

“If people rely too much on agents from the start, they might keep depending on them [...] the conversation might start to feel too standardized. [...] personal style or individuality might shrink.”

Others were uneasy with the agent acting as a “universal adapter” that continually adjusted them to fit a partner’s preferences. P14 warned that *“once you take the agent away, you can’t fit together anymore,”* suggesting that smoothing interaction could paradoxically weaken users’ ability to sustain the relationship independently.

Underlying these concerns was a deeper issue of authorship and responsibility. For participants, the decisions about what to say needed to remain theirs. As P18 stated, *“It should still be my decision [...] it has to be me making the decision and taking responsibility for what I say.”* When agents evaluated, corrected, or generated too much, disclosure risked becoming compliance rather than self-expression, blurring responsibility for what was ultimately disclosed.

Together, these accounts surface a central tension around Disclosure Agency. Participants wanted AI to scaffold reflection, pacing, and learning while insisting that authorship, responsibility, and decision-making remain firmly their own. The value of AI support hinged not on how much guidance it could provide, but on whether it preserved the user’s sense of being the one who decides, speaks, and stands behind what is disclosed.

This tension points to an important design constraint for AI-mediated self-disclosure: agents should support awareness and learning without substituting judgment or action. When systems move from scaffolding into automation (e.g. by deciding what to say, when to say it, or how to phrase it), they risk undermining authorship and responsibility—precisely what users insist on retaining.

4.2 Interpretation: Making Sense of the Partner vs Prematurely Defining the Partner

Early-stage dating conversations are often marked by ambiguity, where minimal or delayed responses make it difficult to know how a message was received or interpreted. Participants discussed moments where they sent disclosures and received little feedback, leaving them unsure whether the issue was timing, topic choice, lack of interest, or external distraction. As P28 explained, *“There are times when I send messages, and the other person barely responds. It’s hard to tell whether it’s because of the topic, their lack of interest, or if they’re simply busy chatting with others.”*

This uncertainty made it difficult for participants to decide how to proceed. They worried about mistaking politeness for engagement, or continuing to elaborate when interest was not actually there. As P11 reflected, *“Every time I think, ‘She’s probably enjoying this conversation,’ she might not be, because I don’t really know what she’s thinking.”* When uncertainty persisted, disclosure could begin to feel burdensome rather than connective. As P23 explained, *“When I share a lot of detail about what I’m doing, but the partner is not really interested in the topic, it becomes a bit of a burden.”*

To navigate these moments, participants envisioned AI agents as helping them interpret ambiguous conversational cues and adjust their disclosures accordingly. They imagined their partner’s agent acting as a lightweight interpersonal translator, offering tentative signals about whether disclosures were likely being received positively.

Participants made it clear that the interpretation needed to be grounded in an understanding of the partner’s conversational style, derived from patterns across interactions. As P12 suggested, *“I think the agent could review the partner’s past chat history to understand how they typically act.”* P07 similarly noted, *“Because his bot knows about him—like his information and stuff—it will be better for his bot to interpret his words for me.”* For instance, knowing that a partner typed slowly or preferred a calmer tone could help them avoid overwhelming the interaction. As P27 explained, *“The agent could analyze that he types more slowly; then I wouldn’t just go on and on talking nonstop.”*

However, participants were also concerned that the same interpretive support that made disclosure feel safer could begin to define the partner too quickly. They emphasized that early dating conversations are often fragmented and surface-level, providing limited data for reliable interpretation. As P20 noted, *“People don’t usually share that much or that fully on a single dating app. If the agent only sees those limited records, it won’t know the full background or reasons.”*

Participants feared that agents might misinterpret subtle cues, such as emojis or habitual expressions. As P16 explained, *“I use ‘==’ to express helplessness, but the agent might read it as annoyance. Everyone has their own habitual ways of expressing things.”* Such misreadings could shift how a partner was perceived and how subsequent disclosures unfolded.

Concerns also extended to how agents might hold onto outdated impressions. Participants emphasized that people change over time, while agents might continue interpreting them through an obsolete lens. As P17 reflected, *“If it just looks at historical data, it might say he’s insecure, but that’s not who he is anymore. For me, it’s much*

better to understand through conversation rather than relying on the agent’s surface analysis.”

These risks felt especially consequential in the early stages of a relationship, when impressions are fragile and easily solidified. Early evaluations might prematurely shape their understanding of a partner. As P18 explained, *“If I’ve only just started getting to know someone, and the agent already gives me an evaluation, I might form a preconceived impression. If it says something negative, I might just stop talking right away.”*

Several participants described a deeper discomfort with seeing their partner through the agent’s perspective rather than forming their own understanding. As P17 noted, *“If the agent interprets the other person too much, then am I seeing them through its perspective instead of getting to know them myself?”* Others likened this effect to how a friend’s casual remark could skew perception. As P22 explained, *“It’s like when a friend says, ‘Oh, he said that because he likes you,’ and you start believing it.”*

For some, the cost of premature interpretation was not just misjudgment, but lost opportunity for discovery. As P14 reflected, *“If the agent told me too early, I’d just label them and stop talking. I’d lose the chance to see their other sides.”*

Together, these accounts surface a central tension around Interpretation. Participants wanted AI support to help them make sense of ambiguous partner cues and avoid misaligned disclosure. At the same time, they worried that interpretive assistance could harden into premature definitions that foreclose gradual discovery.

This tension points to an important design constraint for AI-mediated interpretation in early dating. Agents should support provisional sensemaking without fixing early impressions of who a partner is. Designing for interpretation in intimate contexts therefore requires preserving ambiguity long enough for people to get to know each other through interaction, rather than replacing that process with early conclusions.

4.3 Relational Buffering: Reducing Social Awkwardness vs Preserving Experiential Value

Early dating interactions often involve face-threatening moments, where expressing discomfort, disinterest, or personal boundaries feels socially risky. In such situations, participants described holding back even when conversations drifted into unwanted or uncomfortable territory. As one participant put it, *“I’m afraid I might hurt the other person’s feelings or come across as rude if I said [it] directly”* (P01), when explaining that they might not to talk about a particular topic.

Delegating communication to an AI agent emerged as a way to manage these moments. Letting the agent intervene made sensitive feedback feel softer and less personally exposing. One participant imagined this clearly: *“The agent could nudge my partner privately that they might be talking too much without mentioning my input”* (P02).

The value of relational buffering to participants was in distancing the user from the message. Participants thought that getting this kind of feedback would be easier to receive when framed as system-generated “idea” rather than being user-initiated. As P01 explained, *“The agent shouldn’t say, ‘Your partner doesn’t want to talk about this.’”*

It should sound like a neutral observation, so my partner wouldn't know it came from me."

Relational buffering was also appealing when broaching sensitive questions. Participants wanted to learn about values, intentions, or life circumstances without appearing overly direct or intentional. As P22 noted, *"Sometimes I want to know more personal things about someone, but if I ask directly, it might seem too intentional."*

In these cases, relational buffering also softened the emotional risk of rejection. If a partner ignored an agent's prompt, participants could interpret it as disengagement from the system rather than a personal dismissal. As P07 described, *"I wouldn't feel like my partner was directly rejecting me."*

At the same time, participants were uneasy about how relational buffering might change the texture of early interaction. They described self-disclosure as meaningful because it unfolds through reciprocal turns and shared discovery, rather than having connections surfaced upfront. As P20 reflected, *"If the agent already told them, then when I say something and they reply, 'Oh, I know,' the conversation just ends."* For several participants, feeling a genuine sense of connection depended on discovering common ground together over time. As P18 put it, *"If the agent disclosed that upfront, it would skip the surprise. These connections should unfold slowly, so they feel meaningful."*

Relational buffering also raised concerns about being absent from the moment where this meaning is negotiated. When agents disclosed sensitive information without them present, participants worried they would lose the chance to clarify or add nuance. In P23's words, *"If the agent blurts it out, I lose the chance to explain. That whole process of exploring together is gone."*

These accounts surface a central tension around Relational Buffering. Delegating communication reduced social awkwardness and absorbed face-threatening moments. Yet speaking through an agent also risked removing users from the very interactions where intimacy, meaning, and mutual understanding are co-constructed.

This tension suggests a key design constraint for AI-mediated relational buffering. Agents may help absorb social risk, but should not replace users as active participants in relational sensemaking. Designing for relational buffering in intimate contexts therefore requires limiting what is delegated, so that users remain present in the unfolding experience of getting to know each other.

5 Discussion

5.1 Beyond Topic Suggestions: Supporting Self-Disclosure as Reasoning Under Uncertainty

Past research on conversational scaffolding in online dating has emphasized topic suggestion—providing prompts to sustain dialogue and reduce awkward silences [29, 54]. In our study, "topic" was treated as inseparable from self-disclosure. When topics are tied to self-disclosure, they carry a different weight, implicating not only mutual interest but also privacy boundaries and impression management. Participants emphasized that even a seemingly simple prompt—such as sharing hobbies or past experiences—was meaningful insofar as it required a judgment of comfort and consideration of how disclosure might shape the partner's perception. This reframing suggests that users expect AI support not only in

suggesting topic, but in helping them reason about what to disclose and when.

This finding extends prior work by highlighting that in online dating, conversational scaffolding is often about navigating the risks and rewards of self-disclosure. Participants emphasized that this balance was highly context-dependent. The same detail might be welcomed if revealed later in a relationship but judged inappropriate—or too personal—if shared too early. While the close link between disclosure depth and relational closeness is well established in social penetration theory [2] and in studies of online interaction [23, 27, 31, 55], our findings further highlight the challenge of translating these dynamics into automated support from an AI agent. Some participants hoped AI agents could alert them when they were revealing "too much, too soon," whereas others stressed that they already had clear personal boundaries about what should be disclosed to whom, and at what pace. For these users, system recommendations were useful only if coupled with strong controls: they wanted to specify what was off-limits and ensure the system respected those boundaries, even if past records suggested otherwise.

Beyond offering disclosure guidance, participants emphasized the need for interpretive support—particularly in making sense of ambiguous language, tone, emojis, and subtle cues in their partner's messages. While prior work has shown that online daters anticipate how their disclosures will be received and adjust their strategies accordingly [55], our findings extend this and suggest that such anticipation often breaks down when participants feel uncertain about how to interpret their partner's reactions. This challenge was especially pronounced in early-stage conversations, where communication norms had not yet been established and cues were often too vague to decode with confidence—echoing concerns from earlier studies [14, 18, 34, 38, 49]. As a result, participants struggled to assess whether their disclosures were well-received, which in turn made it difficult to decide whether to deepen or slow the pace of sharing. To address this gap, many envisioned the agent not just as a messenger, but as an interpretive assistant—able to help decode the partner's communication style and offer timely guidance to support more attuned and calibrated disclosure.

These expectations were accompanied by concerns about the limits of AI knowledge. Participants noted that agents might misinterpret subtle cues or rely on incomplete or outdated conversational histories, particularly in early interactions where little context is available. Such limitations raised doubts about the reliability of agent interpretations and concerns that suggestions might appear more certain than warranted.

Participants also worried about the timing of AI supports: disclosures or interpretations offered too early could fix impressions before partners had a chance to discover each other through interaction. Even when factually correct, disclosures or interpretations offered too early were seen as preventing compatibility from emerging through interaction. This concern resonates with prior work showing that users are cautious about AI systems making socially consequential interpretations too early, particularly when those interpretations carry the weight of personalization [11].

Participants also emphasized that they did not want to remain perpetually dependent on AI for conversational success. Instead, they hoped to learn from the agent—to identify which disclosures

resonated, which felt premature, and what patterns characterized effective interaction. This reflects a desire to move from short-term assistance toward long-term growth, positioning the agent as a coach. Participants imagined analytic feedback that would help them reflect on their own behaviors, increase self-awareness, and eventually gain independence in managing disclosure.

Our findings suggest that agents should act as restrained scaffolds that support user reasoning under uncertainty. For instance, they could selectively withhold or qualify responses, making uncertainty visible, and avoiding overconfident interpretations based on partial information. By doing so, agents can reduce over-reliance and support users in developing their own capacity to navigate self-disclosure.

5.2 Preserving User Agency in AI-Mediated Dating Interactions

While participants welcomed AI support for self-disclosure, they also raised concerns about how such support might reshape the dynamics of early dating interactions, where impressions form quickly and mutual understanding emerges through interaction.

A central concern was that agent involvement could shape how impressions are formed. Participants emphasized that early-stage conversations should allow impressions to develop gradually through interaction, rather than being influenced by system-generated suggestions or interpretations.

Participants were particularly wary of relying on agents to mediate how they perceive or respond to a partner. When agents take a more active role in suggesting responses or framing interpretations, they risk shifting the process of sensemaking away from the user. This raised concerns that the agent could begin to function as a participant in the relationship rather than a supporting tool.

Participants also expressed concern that AI involvement could subtly alter how authentic an interaction felt. Even when agent contributions were effective, their presence raised questions about authorship and credibility—whether messages still reflected the person behind them or an optimized version shaped by the system—echoing prior concerns that perceived AI involvement can undermine trust and authenticity in dating and online interactions [3, 26, 50]. Over time, participants worried that such guidance could steer them toward a generic or “ideal” way of speaking, gradually flattening personal voice and reducing the distinctiveness of self-presentation [1]. Hence, participants were cautious about granting agents strong authority in shaping how they expressed themselves, even when assistance was framed as supportive rather than substitutive.

Relatedly, while participants saw value in delegating certain sensitive conversational moves—such as softening feedback or introducing difficult topics through a third-party tone—these forms of mediation introduced tension. By offloading interactional work to the agent, users risked becoming less directly engaged in the process of building intimacy.

Thus, we recommend considering agents as supportive assistants rather than actors that shape the relationship. In particular, we recommend avoiding interventions that steer relational outcomes, and limiting the agent’s role in shaping conversational direction. This

helps ensure that users remain the primary participants in the interaction. This helps to preserve user control over how impressions are formed, how disclosures unfold, and how communication is authored. By maintaining this boundary, designers can preserve authenticity and user ownership in early-stage dating interactions.

5.3 Toward Multi-Agent Support for Self-Disclosure in Online Dating

In our interviews, we did not specify whether the agent should belong to the dating platform or to the individual user. We asked participants what kind of support they would want from an agent during early interactions. Yet participants themselves consistently referred to “my agent” or “their agent,” implicitly treating agents as personal assistants rather than platform-level services. This framing extended naturally to the idea that both sides in an interaction would have their own agents—introducing not only a dyadic interaction between two people but also potential exchanges between “my agent” and “their agent.”

Through this framing, participants highlighted that many of the supports they envisioned could not be achieved by a single user-owned agent alone. For example, deciding what to disclose required both self-knowledge and insight into the partner’s likely interests. While a user’s own agent could help manage boundaries around what was appropriate to share, only the partner’s agent—drawing on that partner’s prior interactions—could anticipate which topics might resonate. Similarly, participants imagined communicating with their own agent to set disclosure rules (“what can and cannot be said”), while also querying the partner’s agent for clarifications about ambiguous expressions or interests.

Taken together, these accounts suggest that future systems may be best conceived not as single-agent models but as ecosystems of multiple interacting agents. In such a system, agents could support users at three distinct temporal points (Figure 2).

First, at a general, long-term level, participants imagined their agent acting as a coach across multiple relationships—helping them refine strategies for pacing disclosure, managing impressions, and improving conversational skills. This included reflecting on past interactions, analyzing which disclosures resonated or felt premature, and establishing default rules about what could generally be delegated.

Second, in the moment of conversation with a specific partner, participants desired more immediate, context-aware support. They imagined agents stepping in with real-time interventions—gently reminding them to pace their disclosures by sharing “not too early, and not all at once,” or offering interpretive assistance when their partner’s messages were ambiguous.

Third, in preparation for a particular dating partner, users could work with their agent outside of real-time conversation. Here, they might specify what the agent could disclose on their behalf to that partner, review patterns in that specific partner’s engagement, and even query the partner’s agent about which topics might sustain meaningful interaction.

However, the introduction of multiple agents may also amplify existing concerns. Participants noted that agents could mislead users, and when both sides rely on their respective agents, such misalignment could result in mutual misunderstanding between




Figure 2: Multi-agent support for self-disclosure in online dating at three temporal points. In long-term coaching, the user’s agent provides guidelines, strategies, and reflection across relationships. In in-conversation support, the user’s agent regulates pacing and flags inappropriate disclosures, while the partner’s agent clarifies ambiguous messages. In partner-specific inquiry, the user’s agent negotiates disclosure boundaries, while the partner’s agent contributes insights into engagement patterns and suggests sustaining topics.

partners. Additionally, as more agent-mediated interactions take place, participants worried that users’ direct involvement in the relationship could diminish. Moments of meaning-making and co-construction—central to developing a relationship—might be increasingly taken over by agents, leading users to question who is truly participating in the interaction.

Together, these visions point toward a multi-agent paradigm for supporting disclosure in online dating—one in which both partners bring their own agents to the interaction, and where agents coordinate across levels and timescales. At the same time, many open questions remain. It is unclear, for example, how comfortable users would be with direct exchanges between agents, or what boundaries should govern cross-agent communication. Future research will need to explore these possibilities through the design and field evaluation of multi-agent prototypes, examining how such systems shape disclosure practices, relational dynamics, and user trust.

6 Limitations

Our study has several limitations that should guide future work. First, we relied on semi-structured interviews with illustrative provocations. While our exploratory approach was appropriate for our research goals, it is not necessarily predictive of how people will really behave with real systems [4]. Future research should therefore validate and extend our findings with prototypes and field deployments that examine how AI mediation unfolds in situ. Second, our recruitment focused on early-stage romantic relationship formation, but people’s disclosure practices and reciprocity norms will likely vary depending on their relationship goals [14, 37]. Understanding how AI mediation might adapt to different relational intentions remains an important direction for future work.

7 Conclusion

Our study shows that users see clear opportunities for AI to scaffold self-disclosure in online dating, while simultaneously surfacing

three central design tensions: *Disclosure Agency*, *Interpretation*, and *Relational Buffering*. Participants welcomed support that paced conversations, clarified partner cues, and reduced social awkwardness, yet insisted that authorship, discovery, and co-construction remain firmly under their control. This reframes disclosure from an individual optimization problem to a relational process unfolding between partners and AI agents. By translating these insights into design, dating platforms—used by millions daily—could better support early conversations that are both guided and exploratory, allowing intimacy to emerge with confidence, curiosity, and user agency.

Acknowledgments

We acknowledge with gratitude the support from the National Science and Technology Council, Taiwan (Grant No. 113-2222-E-A49-009-MY3). We also extend our sincere appreciation to all study participants, anonymous reviewers, and our colleagues for their insightful feedback and support, which significantly informed and strengthened this work.

References

- [1] Dhruv Agarwal, Mor Naaman, and Aditya Vashistha. 2025. AI suggestions homogenize writing toward western styles and diminish cultural nuances. In *Proceedings of the 2025 CHI Conference on Human Factors in Computing Systems*. 1–21.
- [2] Irwin Altman and Dalmas A Taylor. 1973. *Social penetration: The development of interpersonal relationships*. Holt, Rinehart & Winston.
- [3] Meryem Barkallah, Yosr Aissa, and Douglas Zytco. 2025. Transparent Hearts: Balancing Privacy and Trust in AI-Generated Self-Presentation for Online Dating. In *Proceedings of the Extended Abstracts of the CHI Conference on Human Factors in Computing Systems*. 1–8.
- [4] Eric PS Baumer, Vera Khovanskaya, Mark Matthews, Lindsay Reynolds, Victoria Schwanda Sosik, and Geri Gay. 2014. Reviewing reflection: on the use of reflection in interactive system design. In *Proceedings of the 2014 conference on Designing interactive systems*. 93–102.
- [5] John H Berg and Margaret S Clark. 1986. Differences in social exchange between intimate and other relationships: Gradually evolving or quickly apparent? In *Friendship and social interaction*. Springer, 101–128.
- [6] Mark Blythe. 2014. Research through design fiction: narrative in real and imaginary abstracts. In *Proceedings of the SIGCHI conference on human factors in*

- computing systems. 703–712.
- [7] Virginia Braun and Victoria Clarke. 2006. Using thematic analysis in psychology. *Qualitative research in psychology* 3, 2 (2006), 77–101.
 - [8] Penelope Brown. 1987. *Politeness: Some universals in language usage*. Vol. 4. Cambridge university press.
 - [9] Jacob L Cayanus and Matthew M Martin. 2008. Teacher self-disclosure: Amount, relevance, and negativity. *Communication Quarterly* 56, 3 (2008), 325–341.
 - [10] Nancy L Collins and Lynn Carol Miller. 1994. Self-disclosure and liking: a meta-analytic review. *Psychological bulletin* 116, 3 (1994), 457.
 - [11] Isha Datey and Douglas Zytco. 2024. "Just Like, Risking Your Life Here": Participatory Design of User Interactions with Risk Detection AI to Prevent Online-to-Offline Harm Through Dating Apps. *Proceedings of the ACM on Human-Computer Interaction* 8, CSCW2 (2024), 1–41.
 - [12] Valerian J. Derlega and Janusz Grzelak. 1979. Appropriateness of self-disclosure. In *Self-disclosure: Origins, patterns, and implications of openness in interpersonal relationships*, Gordon J. Chelune (Ed.), Jossey-Bass, 151–176.
 - [13] Valerian J Derlega, Sandra Metts, Sandra Petronio, and Stephen T Margulis. 1993. *Self-disclosure*. Sage Publications, Inc.
 - [14] Nicole Ellison, Rebecca Heino, and Jennifer Gibbs. 2006. Managing impressions online: Self-presentation processes in the online dating environment. *Journal of computer-mediated communication* 11, 2 (2006), 415–441.
 - [15] Yue Fu, Yixin Chen, Zelia Gomes Da Costa Lai, and Alexis Hiniker. 2024. Should ChatGPT write your breakup text? Exploring the role of AI in relationship dissolution. *arXiv preprint arXiv:2401.09695* (2024).
 - [16] Bill Gaver, Tony Dunne, and Elena Pacenti. 1999. Design: cultural probes. *interactions* 6, 1 (1999), 21–29.
 - [17] Jennifer L Gibbs, Nicole B Ellison, and Rebecca D Heino. 2006. Self-presentation in online personals: The role of anticipated future interaction, self-disclosure, and perceived success in Internet dating. *Communication research* 33, 2 (2006), 152–177.
 - [18] Jennifer L Gibbs, Nicole B Ellison, and Chih-Hui Lai. 2011. First comes love, then comes Google: An investigation of uncertainty reduction strategies and self-disclosure in online dating. *Communication Research* 38, 1 (2011), 70–100.
 - [19] Thalia R Goldstein and Ellen Winner. 2012. Enhancing empathy and theory of mind. *Journal of cognition and development* 13, 1 (2012), 19–37.
 - [20] Kathryn Greene, Valerian J Derlega, and Alicia Mathews. 2006. Self-disclosure in personal relationships. *The Cambridge handbook of personal relationships* 409 (2006), 427.
 - [21] Jeffrey T Hancock, Mor Naaman, and Karen Levy. 2020. AI-mediated communication: Definition, research agenda, and ethical considerations. *Journal of Computer-Mediated Communication* 25, 1 (2020), 89–100.
 - [22] Jennifer R Henretty and Heidi M Levitt. 2010. The role of therapist self-disclosure in psychotherapy: A qualitative review. *Clinical psychology review* 30, 1 (2010), 63–77.
 - [23] Lee Bee Hian, Sim Li Chuan, Tan Mon Kiat Trevor, and Benjamin H Detenber. 2004. Getting to know you: Exploring the development of relational intimacy in computer-mediated communication. *Journal of computer-mediated communication* 9, 3 (2004), JCMC935.
 - [24] Clara E Hill and Sarah Knox. 2001. Self-disclosure. *Psychotherapy: Theory, Research, Practice, Training* 38, 4 (2001), 413.
 - [25] Jess Hohenstein and Malte Jung. 2020. AI as a moral crumple zone: The effects of AI-mediated communication on attribution and trust. *Computers in Human Behavior* 106 (2020), 106190.
 - [26] Maurice Jakesch, Megan French, Xiao Ma, Jeffrey T Hancock, and Mor Naaman. 2019. AI-mediated communication: How the perception that profile text was written by AI affects trustworthiness. In *Proceedings of the 2019 CHI conference on human factors in computing systems*. 1–13.
 - [27] L Crystal Jiang, Natalie N Bazarova, and Jeffrey T Hancock. 2011. The disclosure-intimacy link in computer-mediated communication: An attributional extension of the hyperpersonal model. *Human communication research* 37, 1 (2011), 58–77.
 - [28] Sidney M Jourard. 1971. Self-disclosure: An experimental analysis of the transparent self. (1971).
 - [29] Taewook Kim, Qingyu Guo, Hyeonjae Kim, Wenjie Yang, Meiziniu Li, and Xiaojuan Ma. 2022. Facilitating Continuous Text Messaging in Online Romantic Encounters by Expanded Keywords Enumeration. In *Companion Publication of the 2022 Conference on Computer Supported Cooperative Work and Social Computing*. 3–7.
 - [30] Jean-Philippe Laurenceau, Lisa Feldman Barrett, and Paula R Pietromonaco. 1998. Intimacy as an interpersonal process: the importance of self-disclosure, partner disclosure, and perceived partner responsiveness in interpersonal exchanges. *Journal of personality and social psychology* 74, 5 (1998), 1238.
 - [31] Juwon Lee, Omri Gillath, and Andrew Miller. 2019. Effects of self- and partner's online disclosure on relationship intimacy and satisfaction. *PloS one* 14, 3 (2019), e0212186.
 - [32] Leah E LeFebvre. 2018. Swiping me off my feet: Explicating relationship initiation on Tinder. *Journal of Social and Personal Relationships* 35, 9 (2018), 1205–1229.
 - [33] Li Li, Peilin Cai, Ryan A Rossi, Franck Dernoncourt, Branislav Kveton, Junda Wu, Tong Yu, Linxin Song, Tiankai Yang, Yuehan Qin, et al. 2025. A Personalized Conversational Benchmark: Towards Simulating Personalized Conversations. *arXiv preprint arXiv:2505.14106* (2025).
 - [34] Louise Anker Nexø and Jacob Strandell. 2020. Testing, filtering, and insinuating: Matching and attunement of emoji use patterns as non-verbal flirting in online dating. *Poetics* 83 (2020), 101477.
 - [35] Pew Research Center. 2023. Key findings about online dating in the U.S. <https://www.pewresearch.org/short-reads/2023/02/02/key-findings-about-online-dating-in-the-u-s/>. Accessed: 2025-09-11.
 - [36] Belle Rose Ragins. 2012. Reflections on the researcher-practitioner gap: A call for professionalization of HR. *Human Resource Management* 51, 5 (2012), 745–759. doi:10.1002/hrm.21494
 - [37] Giulia Ranzini and Christoph Lutz. 2017. Love at first swipe? Explaining Tinder self-presentation and motives. *Mobile Media & Communication* 5, 1 (2017), 80–101.
 - [38] David L Rodrigues, Bernardo P Cavalheiro, and Marília Prada. 2022. Emoji as Icebreakers? Emoji can signal distinct intentions in first time online interactions. *Telematics and Informatics* 69 (2022), 101783.
 - [39] Larry D Rosen, Nancy A Cheever, Cheyenne Cummings, and Julie Felt. 2008. The impact of emotionality and self-disclosure on online dating versus traditional dating. *Computers in Human Behavior* 24, 5 (2008), 2124–2157.
 - [40] Michael J Rosenfeld, Reuben J Thomas, and Sonia Hausen. 2019. Disintermediating your friends: How online dating in the United States displaces other ways of meeting. *Proceedings of the National Academy of Sciences* 116, 36 (2019), 17753–17758.
 - [41] Novita Sari Satyaningrum et al. 2025. Social Penetration in Digital Spaces: Developing Interpersonal Relationships Through the Bumble Dating App. *International Journal of Research and Innovation in Social Science* 9, 3 (2025), 2220–2229.
 - [42] Joanna Schug, Masaki Yuki, and William Maddux. 2010. Relational mobility explains between-and within-culture differences in self-disclosure to close friends. *Psychological Science* 21, 10 (2010), 1471–1478.
 - [43] Ashish Sharma, Inna W Lin, Adam S Miner, David C Atkins, and Tim Althoff. 2023. Human-AI collaboration enables more empathic conversations in text-based peer-to-peer mental health support. *Nature Machine Intelligence* 5, 1 (2023), 46–57.
 - [44] R d Shaver and P Reis. 1988. *Intimacy as an interpersonal process*. USA.
 - [45] Donghoo Shin, Soomin Kim, Ruoxi Shang, Joonhwan Lee, and Gary Hsieh. 2023. IntroBot: Exploring the use of chatbot-assisted familiarization in online collaborative groups. In *Proceedings of the 2023 CHI Conference on Human Factors in Computing Systems*. 1–13.
 - [46] Patricia M Sias. 2005. Workplace relationship quality and employee information experiences. *Communication studies* 56, 4 (2005), 375–395.
 - [47] Susan Sprecher, Stanislav Treger, Joshua D Wondra, Nicole Hilaire, and Kevin Walpe. 2013. Taking turns: Reciprocal self-disclosure promotes liking in initial interactions. *Journal of Experimental Social Psychology* 49, 5 (2013), 860–866.
 - [48] Nu Tang, Lisamarie Bensman, and Elaine Hatfield. 2013. Culture and sexual self-disclosure in intimate relationships. *Interpersona: An International Journal on Personal Relationships* 7, 2 (2013), 227–245.
 - [49] Joseph B. Walther. 1996. Computer-mediated communication: Impersonal, interpersonal, and hyperpersonal interaction. *Communication Research* 23, 1 (1996), 3–43. doi:10.1177/009365096023001001
 - [50] Yihan Wu and Ryan M Kelly. 2020. Online dating meets artificial intelligence: How the perception of algorithmically generated profile text impacts attractiveness and trust. In *Proceedings of the 32nd Australian conference on human-computer interaction*. 444–453.
 - [51] Fan Zhang, Yun Chen, Xiaoke Zeng, Tianqi Wang, Long Ling, and Ray Lc. 2025. "An Image of Ourselves in Our Minds": How College-educated Online Dating Users Construct Profiles for Effective Self-Presentation. *Proceedings of the ACM on Human-Computer Interaction* 9, 2 (2025), 1–30.
 - [52] Li Zhou, Jianfeng Gao, Di Li, and Heung-Yeung Shum. 2020. The design and implementation of xiaoice, an empathetic social chatbot. *Computational Linguistics* 46, 1 (2020), 53–93.
 - [53] Douglas Zytco and Nicholas Furlo. 2023. Online dating as context to design sexual consent technology with women and lgbtq+ stakeholders. In *Proceedings of the 2023 CHI conference on human factors in computing systems*. 1–17.
 - [54] Doug Zytco, Sukeshini Grandhi, and Quentin Jones. 2016. The coaches said... what? Analysis of online dating strategies recommended by dating coaches. In *Proceedings of the 2016 ACM International Conference on Supporting Group Work*. 385–397.
 - [55] Douglas Zytco, Sukeshini A Grandhi, and Quentin Jones. 2014. Impression management struggles in online dating. In *Proceedings of the 2014 ACM International Conference on Supporting Group Work*. 53–62.
 - [56] Douglas Zytco, Devin Tebbe, Braeden Burger, Emma Walquist, and Meryem Barkallah. 2025. Beyond Affirmative Consent: Creating Alternative Models for Consent to Computer-Mediated Sexual Activity in the Context of VR Dating. (2025).

A AI Mediation Provocations in a Reciprocity Scenario




Figure 3: AI mediation provocations used as design probes, illustrated through a reciprocity scenario in which the participant is contributing more to the conversation than their partner. Provocations were organized along two dimensions: *Perspective* and *Delivery Mode*. Under *Perspective*, prompts either focused on *User Behavior*, highlighting what the participant had done in the interaction (e.g., “You’ve been contributing more to the conversation than your partner”), or on *Partner Experience*, highlighting how the partner might experience the interaction (e.g., “Your partner may find it difficult to find space to contribute to the conversation”). Under *Delivery Mode*, prompts were either *Public*, visible to both the participant and partner, or *Private*, visible only to the participant.